

# The Last Mile to CAL Compiler for Epiphany Architecture

Mingkun Yang

HH

September 2, 2013

# Outline

- 1 Introduction
- 2 Method
  - Incremental Refinement
  - Asynchronous and Synchronous Call
  - Synchronization
  - Manage all actors
- 3 Result
  - IDCT
  - Tiny buffer size
  - Medium buffer size
  - Large buffer size
- 4 Conclusion
- 5 Bonus: two mysterious bugs

# Table of Contents

- 1 Introduction
- 2 Method
  - Incremental Refinement
  - Asynchronous and Synchronous Call
  - Synchronization
  - Manage all actors
- 3 Result
  - IDCT
  - Tiny buffer size
  - Medium buffer size
  - Large buffer size
- 4 Conclusion
- 5 Bonus: two mysterious bugs

# Introduction

- Multiple-core trend
  - The end of higher frequency
  - Scale horizontally
- CAL Actor Language
  - DSL of dataflow program
  - Explicit parallelism

# Brief intro to CAL

- Actor

## Listing 1 : Identity Actor

```
actor ID () In ==> Out :  
    action In: [a] ==> Out: [a] end  
end
```

---

- Network

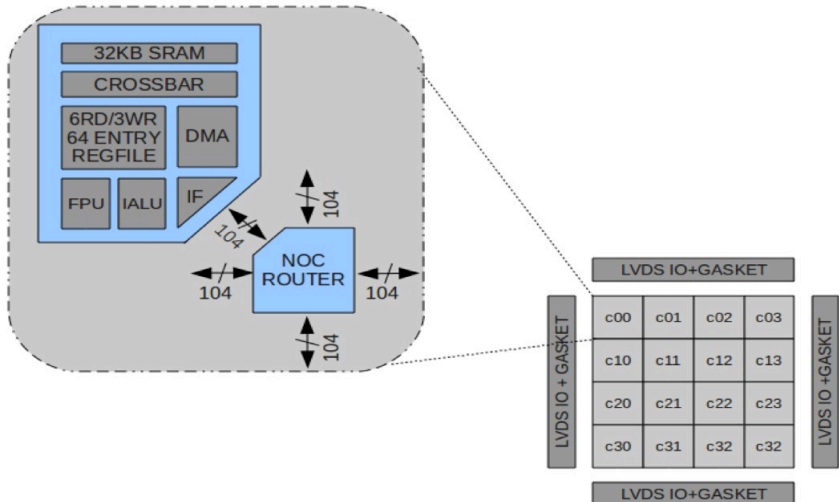
## Listing 2 : Simple network

```
x = ID();  
y = ID();  
...  
x.out --> y.in;  
...
```

## Related work

- OpenDF
  - ACTORS project (d2c)[1]
- Orcc
  - Synthesis from Dataflow Programs[2]

# Epiphany Architecture



# Table of Contents

## 1 Introduction

## 2 Method

- Incremental Refinement
- Asynchronous and Synchronous Call
- Synchronization
- Manage all actors

## 3 Result

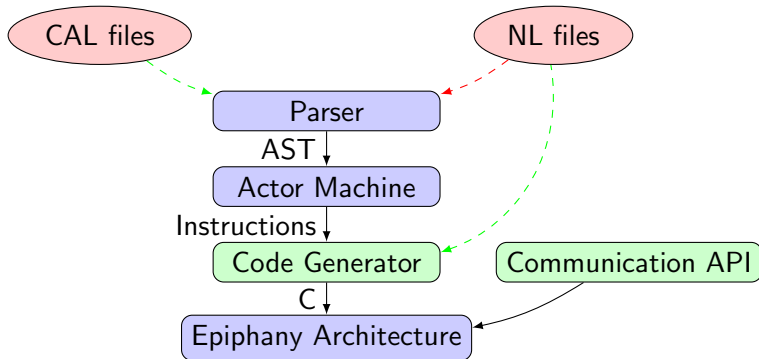
- IDCT
- Tiny buffer size
- Medium buffer size
- Large buffer size

## 4 Conclusion

## 5 Bonus: two mysterious bugs



# CAL compiling process




# Code Generator

- Source-to-source compiler
  - Readability
  - Actor-based translation
    - One-to-one mapping (from .cal to .c)
    - Structure for one actor (Classes in OOP)
  - No duplication in linking
- Build Process
  - Eclipse (Default)
    - Each project for each core.
  - Customized Makefile
    - Stop on any failure from any core.


# Communication API

- `epiphany_write`
- `epiphany_read`
- `port_end`
  - Called when there's no further tokens in this transaction.
  - Resembles end of packet in communication protocol.
  - Better measurement of active duration of each actor (each core).
- `connect`
  - Connect one input port with one output port.

# Destination-buffer




## Destination-buffer




- Local > Remote

# Both-buffer




## Both-buffer



- Direct Memory Access (DMA)
  - Separate “thread”

# Double-buffer




- Will be blocked if one is faster than the other




# Design and Examples

- Design
  - Async by default (Never block)
  - Only block when necessary (Fall back to sync call)
  - Push tokens to the destination core(s).
- Example
  - try\_flush vs do\_flush
  - try\_distribute vs do\_distribute


## try\_flush flow



(a) Initial condition




(b) Write to the last slot




(c) Index pointer is updated

# do\_flush


- Red: unknown or uninterested
- Blue: occupied




# try\_distribute



(a) One fifo becomes full




(b) Write when one fifo is full

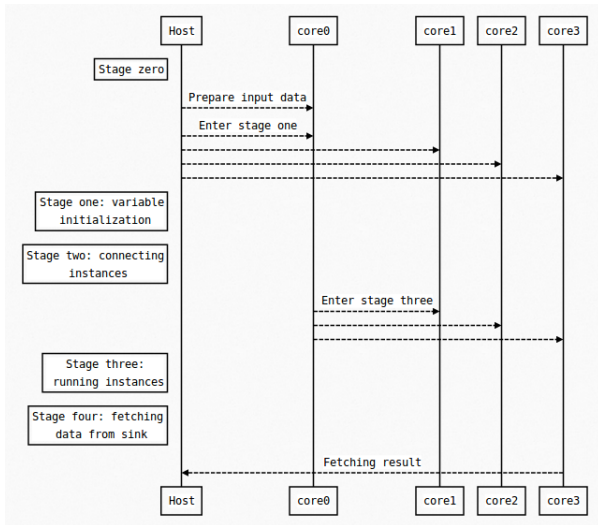


(c) Another write operation

# do\_distribute



# Synchronization between the board and host



# Thread like management

- not\_finished
  - There are further tokens from any of input ports.
  - Or there are tokens in the local buffer.
- run
  - Fire actions, consume and produce tokens.
- end
  - Process all tokens in the local buffer.
  - Mark the end-of-token in this port.

# Strategy Pattern

## Listing 3 : How Actor Interface is Used

```
// common/common.c
...
void core_main(void *a, init_t *init) {
    ...
    while(api->not_finished(a)) {
        api->run(a);
    }
    api->end(a);
    ...
}
...
```

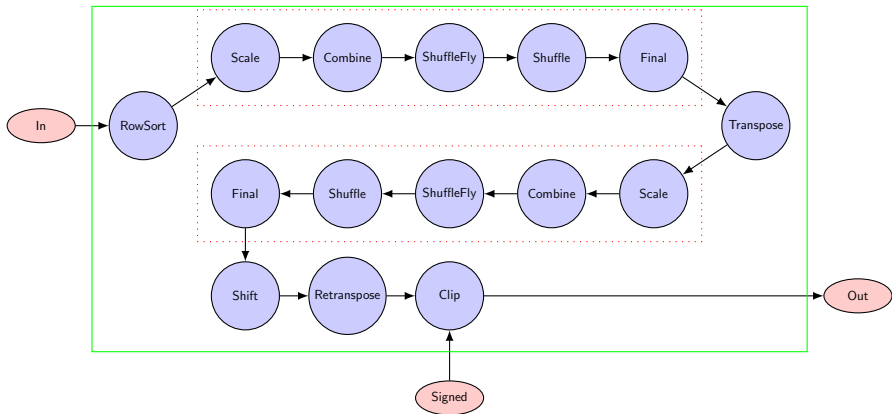
---



# Table of Contents


- 1 Introduction
- 2 Method
  - Incremental Refinement
  - Asynchronous and Synchronous Call
  - Synchronization
  - Manage all actors
- 3 **Result**
  - IDCT
  - Tiny buffer size
  - Medium buffer size
  - Large buffer size
- 4 Conclusion
- 5 Bonus: two mysterious bugs

# IDCT2D




- 64000 tokens for *In* and 1000 tokens for *Signed*
- 100 times, and get the average


# Tiny buffer size: api call vs total execution




(a) destination buffer version



(b) both buffer version




(c) double buffer version




(d) maximum from all three version


# Tiny buffer size: polling vs api call




(a) destination buffer version



(b) both buffer version




(c) double buffer version




(d) maximum from all three version


# Medium buffer size: api call vs total execution




(a) destination buffer version



(b) both buffer version




(c) double buffer version




(d) maximum from all three version


# Medium buffer size: polling vs api call




(a) destination buffer version



(b) both buffer version




(c) double buffer version




(d) maximum from all three version


# Large buffer size: api call vs total execution




(a) destination buffer version



(b) both buffer version




(c) double buffer version




(d) maximum from all three version


# Large buffer size: polling vs api call




(a) destination buffer version



(b) both buffer version



(c) double buffer version



(d) maximum from all three version



# Table of Contents

- 1 Introduction
- 2 Method
  - Incremental Refinement
  - Asynchronous and Synchronous Call
  - Synchronization
  - Manage all actors
- 3 Result
  - IDCT
  - Tiny buffer size
  - Medium buffer size
  - Large buffer size
- 4 Conclusion
- 5 Bonus: two mysterious bugs

# On the Shoulder of Giants

- Front end
- Actor Machine
- Sequential C Code Generator
- The communication API for Epiphany

## Conclusion and future work

- Actor is lightweight, and it's better to assign multiple actors to one core.
  - Actor composition in Actor Machine
  - multiple actors in one core – user space thread

## Conclusion and future work

- Actor is lightweight, and it's better to assign multiple actors to one core.
  - Actor composition in Actor Machine
  - multiple actors in one core – user space thread
- Balance the effort between creating tests and extra caution from developers
  - Create reasonable amount of unit tests.
  - Be focused while developing.

# Resource

 <http://www.actors-project.eu/>



Ghislain Roquier

Hardware and Software Synthesis of Heterogeneous Systems  
from Dataflow Programs

 <http://www.bdti.com/InsideDSP/2012/09/05/Adapteva>

 <https://bitbucket.org/albertnetymk/epiphany>

# Table of Contents

- 1 Introduction
- 2 Method
  - Incremental Refinement
  - Asynchronous and Synchronous Call
  - Synchronization
  - Manage all actors
- 3 Result
  - IDCT
  - Tiny buffer size
  - Medium buffer size
  - Large buffer size
- 4 Conclusion
- 5 Bonus: two mysterious bugs

## Listing 4 : buggy code

```
typedef struct {  
    int players;  
    int source[20];  
    ...  
}
```

## Listing 5 : fix

```
typedef struct {  
    int dummy;  
    int players;  
    int source[20];  
    ...  
}
```

## Listing 4 : buggy code

```
typedef struct {  
    int players;  
    int source[20];  
    ...  
}
```

## Listing 5 : fix

```
typedef struct {  
    int dummy;  
    int players;  
    int source[20];  
    ...  
}
```

- Observation



## Listing 4 : buggy code

```
typedef struct {  
    int players;  
    int source[20];  
    ...  
}
```

## Listing 5 : fix

```
typedef struct {  
    int dummy;  
    int players;  
    int source[20];  
    ...  
}
```

- Observation
  - Consistent and deterministic.

## Listing 4 : buggy code

```
typedef struct {  
    int players;  
    int source[20];  
    ...  
}
```

## Listing 5 : fix

```
typedef struct {  
    int dummy;  
    int players;  
    int source[20];  
    ...  
}
```

- Observation
  - Consistent and deterministic.
  - 64 bit host

## Listing 4 : buggy code

```
typedef struct {  
    int players;  
    int source[20];  
    ...  
}
```

## Listing 5 : fix

```
typedef struct {  
    int dummy;  
    int players;  
    int source[20];  
    ...  
}
```

- Observation
  - Consistent and deterministic.
  - 64 bit host
  - `sizeof(int) == 4`

## Listing 4 : buggy code

```
typedef struct {  
    int players;  
    int source[20];  
    ...  
}
```

## Listing 5 : fix

```
typedef struct {  
    int dummy;  
    int players;  
    int source[20];  
    ...  
}
```

- Observation
  - Consistent and deterministic.
  - 64 bit host
  - `sizeof(int) == 4`
- Explanation
  - field alignment in structure

## Listing 6 : buggy code

```
bool might_has_input(port_in *p)
{
    return has_input(p, 1) || !p->end ;
}
```

---

## Listing 7 : fix

```
bool might_has_input(port_in *p)
{
    return !p->end || has_input(p, 1) ;
}
```

---

- Observation

- Observation
  - Non-deterministic.

- Observation
  - Non-deterministic.
  - precedence, assembly, ...



- Observation
  - Non-deterministic.
  - precedence, assembly, ...
- Explanation (with external help)
  - Shared resource is modified between two checks.
  - Check them in the reverse order of mutating.